

PaperTrail: A Retrieval Augmented System for Automatic Dataset Description Generation

Maria Beatriz Silva
ms14127@nyu.edu

Sara Vargas Martínez
sv2279@nyu.edu

Abstract

Automated dataset description approaches rely on statistical and semantic profiling of the data itself, missing crucial contextual information about collection methodology, intended use cases, and known limitations that may appear in citing literature. We present PaperTrail, an extension to the AutoDDG (Zhang et al., 2025) dataset description system, that incorporates retrieval-augmented generation to enhance descriptions with insights from scientific papers that cite the dataset. Our method adds a module to AutoDDG’s pipeline that: identifies relevant sections in citing papers and extracts information about the dataset from these sections. This enables research context to be integrated with AutoDDG’s statistical and semantic profiles. In a controlled experiment across 10 scientific datasets from Zenodo, PaperTrail improves the coverage, as defined in 5, of a generated description by 6.2% over the baseline AutoDDG system over all profile extraction prompts tested. These improvements demonstrate that research literature contains complementary information to statistical profiles, enabling more comprehensive dataset descriptions that better serve researchers’ needs for understanding dataset provenance, usage patterns, and scientific context. We release our code as an open-source extension to AutoDDG.

1 Introduction

Scientific datasets are more accessible than ever. Repositories like Zenodo, Hugging Face, and domain-specific archives host millions of datasets spanning genomics, social science, natural language processing, and beyond. Yet despite this abundance, finding relevant datasets remains difficult due to insufficient or non-existent dataset descriptions. Automated dataset description systems offer a partial solution to this discovery problem. Systems like AutoDDG (Zhang et al., 2025) generate natural language summaries by profiling data

statistics—distributions, types, missing values—and prompting large language models (LLMs) to derive semantic information (e.g., the dataset topic, semantic types of attributes); this information is then fed to an LLM to create a dataset description. These detailed descriptions can be indexed by dataset search systems, improving findability, and they also help researchers assess dataset relevance by providing readable and concise information about dataset’s structure and content. However, they miss critical information about how researchers can actually *use* the dataset: what analyses it supports well, what biases to account for, what pre-processing steps prior work found necessary, and what limitations of the data may be. This information may be available online, but not in the datasets themselves.

We explore the use of scientific literature as an additional source of context information for datasets. When researchers publish work using a dataset, they often document aspects of their experience, such as methodological decisions in response to data characteristics, what problems the data allowed them to solve, pre-processing strategies they employed, as well as the data provenance.

We propose PaperTrail, a system that, given a dataset and a paper that cites it, extracts dataset-relevant information from the paper and integrates this research context alongside intrinsic information from the dataset itself (see Section 3). Specifically, PaperTrail identifies relevant sections in the paper extracts information about collection methodology, usage patterns, and limitations—the practical knowledge researchers need for informed dataset selection.

A core technical challenge lies identifying and extracting such information from scientific publications, due to the variability of how papers refer to and discuss datasets. Relevant information about datasets may appear in methods, re-

sults, or limitations sections; it may be explicit ("this dataset exhibits selection bias") or implicit (reflected in preprocessing choices); and it must be distinguished from general methodological content unrelated to the specific dataset. We explore different extraction strategies to address these challenges.

We evaluate PaperTrail on 10 dataset-paper pairs from Zenodo, we find 6.2% coverage score improvement over baseline AutoDDG over all prompts. Our results demonstrate that scientific literature contains complementary information to statistical profiles, and that this information can be systematically extracted to produce more comprehensive, usage-oriented dataset descriptions. This project contributes (1) open-source code for PaperTrail and its integration with AutoDDG, (2) curated dataset-paper evaluation pairs with their original dataset descriptions written by the dataset authors for comparison, and (3) an interactive viewer that allows users to compare the descriptions generated by the Vanilla AutoDDG and those with the PaperTrail extension.

2 Related Work

2.1 Automated Dataset Description

AutoDDG (Zhang et al., 2025) proposes a three-stage framework for automated dataset description generation tailored to tabular data. First, it constructs a *data profile* capturing essential statistical characteristics using traditional data profiling techniques. Second, it generates a *semantic profile* through LLM prompting to infer contextual information such as domain and potential use cases beyond what statistics alone reveal. Finally, a Description Generation Engine produces two descriptions: one optimized for human readability and another for dataset retrieval systems.

The authors demonstrate that LLM-generated descriptions significantly improve dataset discoverability, and evaluate their descriptions across dimensions: retrieval performance, reference-based metrics (ROUGE (Ganesan, 2018), BERTScore (Zhang et al., 2020)), and reference-free LLM-based quality assessment. However, AutoDDG has a critical limitation: it relies solely on what can be deduced from the dataset itself without considering how the dataset has been used in practice. While the semantic profile adds context through LLM prompting, this represents general model knowledge rather than dataset-specific usage pat-

terns, limitations, and methodological considerations documented in citing literature. Our work addresses this gap by augmenting AutoDDG’s data-driven and semantic profiles with specific contextual information extracted from scientific papers.

2.2 Retrieval-Augmented Generation

Retrieval-Augmented Generation (RAG) has emerged as a powerful way to enhance language model outputs with external knowledge. The standard RAG approach involves retrieving relevant documents from a corpus using dense retrieval methods, then conditioning a language model on both the query and retrieved context to generate informed responses (Lewis et al., 2021). This pattern has proven particularly effective for tasks requiring up-to-date or domain-specific information beyond a model’s training data (Gao et al., 2024).

While PaperTrail is inspired by RAG’s core insight—external documents can provide valuable context for generation—our approach differs from traditional RAG implementations in several ways. Rather than using embedding-based retrieval across a large corpus, we work with known dataset-paper pairs where the relevant document is specified in advance. Our focus shifts from retrieval to extraction: identifying which sections of a given paper discuss the dataset and extracting relevant information from those sections (see Section 3.1 for details).

This extraction-focused approach is necessitated by the unpredictable ways scientific papers discuss datasets. A dataset may be the paper’s primary focus or merely one of several used in experiments; it may be discussed extensively in methods or mentioned briefly in results; relevant information about limitations may appear in discussion sections or not at all. Our work thus extends the RAG paradigm to a scenario where the challenge is not finding the right document, but finding and extracting the right information from a known document where the dataset’s role, prominence, and treatment are not known in advance.

3 The PaperTrail Module

Our approach augments AutoDDG’s description generation pipeline by adding a PaperTrail module that extracts contextual information from scientific papers citing the dataset, and integrates this information into AutoDDG’s description genera-

tion engine to produce context-enriched descriptions. The pipeline operates in three stages: (1) **PDF Processing**, where we extract text from citing papers and identify sections that discuss or refer to the dataset, (2) **Profile Extraction**, where we use an LLM to extract structured information about the dataset’s methodology, limitations, and usage context from these sections, and (3) **Integration**, where we incorporate this research-derived profile alongside AutoDDG’s existing statistical and semantic profiles in the description generation prompt. Figure 1 illustrates this pipeline.

3.1 PDF Processing & Section Identification

Processing entire papers through an LLM is both expensive and introduces noise that degrades extraction quality. We explored three strategies for identifying dataset-relevant content.

Full Paper Processing: We initially processed entire papers but found that, beyond being expensive, this method overwhelmed the resulting PaperTrail with information about the research paper itself rather than information about the dataset.

Chunked Extraction Strategies: Inspired by RAG approaches, we shifted to chunking papers and sending only relevant chunks. We extract text using PyPDF and partition it into overlapping chunks of 2,000 characters with 200-character overlap to preserve context across boundaries.

We hypothesized that meaningful discussion of dataset properties typically occurs at or near explicit mentions of the dataset itself, suggesting a keyword-based approach might be more effective than semantic similarity. We explored two methods:

- **LLM-based Relevance Scoring:** We prompted Llama 3.1 8B with a chunk and the dataset’s name to classify each chunk as relevant or irrelevant to the dataset. Relevant chunks were then merged and deduplicated.
- **Keyword-based Search:** We implement a token-based anchor search that identifies chunks explicitly referencing the dataset. The process is as follows:
 1. *Anchor Identification:* We tokenize the dataset name, filter common stop words, and search for chunks containing at least 2 key tokens using case-insensitive word-boundary matching. For example, “FLUPrint COVID Dataset” be-

comes key tokens {*FLUPrint*, *COVID*, *Dataset*}, and a relevant chunk will have at least 2 of these tokens (unless there are fewer than 2 key tokens, then 1 will match suffices). We refer to these chunks as *anchor chunks*.

2. *Context Window Expansion:* For each anchor chunk, we extract a context window of ± 3 surrounding chunks. When anchor chunks occur within 6 chunks of each other, we merge them into a single logical block to preserve narrative coherence. Finally, we concatenate all of the relevant chunks to form our full context.
3. *Overlap Resolution:* Since chunks have 200-character overlap and context windows may overlap, we perform character-level deduplication by tracking chunk boundaries and removing redundant content when merging blocks.

We evaluated both strategies on 10 datasets by using the extracted contexts to generate PaperTrails and final descriptions, and then measuring coverage scores. Intuitively, we want to assess if and how additional context contributes to the derived description; for details on the metrics we used, see Section 4.4. As shown in Figure 2, the keyword-based approach achieved higher overall coverage (0.51 vs 0.45) while being more concise—extracting only 78.5% of the paper length compared to 94.1% for LLM-based scoring. The keyword approach particularly excelled at capturing basic information (0.95 vs 0.85) and provenance details (0.33 vs 0.18), suggesting better precision with less noise. Based on these results, we adopted the keyword-based approach for identifying relevant paper sections.

3.2 Profile Extraction

Once relevant paper sections are identified, we extract a structured PaperTrail using an LLM (Llama 3.1 8B). We concatenate all selected chunks into a single text block and prompt the model to extract dataset-specific information from that block.

We explored three prompt design strategies: (1) **Paragraph Summaries** which ask for natural language paragraphs covering research domain, usage patterns, characteristics, and limitations; (2) **Structured Fields** which request a hierarchical JSON schema with predefined fields (i.e., domain, usage, characteristics, and limitations) and explicit




Figure 1: Extended AutoDDG Pipeline with PaperTrail. Our contribution (shown in pink) adds a PaperTrail extracted from citing papers, which is integrated alongside AutoDDG’s existing Data and Semantic Profiles.




Figure 2: Impact of Keyword vs LLM-based Chunk Selection on Description Quality

null handling to ensure the LLM extracts only explicitly stated information; and (3) **Research Decision Support** which frame context extraction around helping researchers evaluate dataset suitability, specifying fields like `enabled_research`, `why_useful`, and `challenges`. We hypothesized that structured outputs might provide more consistent extractions than free-form text, while decision-oriented framing might better align the extracted information with researchers’ practical needs.

We tested multiple variations within each category.¹ See Section 5 for the performance of each prompt.

3.3 Integration with AutoDDG

As shown in Figure 1, AutoDDG’s pipeline generates two core profiles: a data profile capturing statistical characteristics and a semantic profile inferring contextual information about columns. It also generates a “topic” for the dataset, that is, a short phrase identifying the dataset’s domain. AutoDDG’s Description Generation Engine uses a single prompt that aggregates these components

with a dataset sample to produce the final description. The prompt sequentially incorporates each information source, with each component contributing specific sentence(s) to enrich different aspects of the description. We augment this prompt by adding our PaperTrail as an additional information source. Specifically, we insert the following instruction after the semantic and topic profiles:

Additionally, research literature provides this summary about the dataset: {context_profile}. Based on the information within this context, please add high-value sentence(s) to the description describing the most important insights provided by the research, such as how the dataset was acquired, what its key variables are, how it was used in the study, and any specific limitations identified by the paper.

This modification allows AutoDDG to naturally incorporate research-derived context alongside its existing statistical and semantic profiles. This prompt is then fed to the LLM, which produces the description.

4 Evaluation

4.1 Experimental Setup

We created a database of 10 datasets and associated papers found in Table 2. To construct this database, we explored biological and environmental datasets using Zenodo and OpenAlex.² Zenodo was well-suited for our purposes because many records include usable dataset descriptions

¹See our repository for the exact prompts tested.

²These repositories provide open-access datasets with descriptive metadata, and OpenAlex enables discovery of papers that cite a given dataset.

and downloadable files, while OpenAlex allowed us to identify papers citing each dataset. During database construction, we encountered several challenges common to the data repository landscape: sparse metadata, inconsistent citation practices, and difficulty distinguishing papers that substantively use a dataset from those that cite it only tangentially.

To address these challenges and ensure both dataset quality and pipeline compatibility, we applied the following selection criteria: the dataset must be (1) in tabular form (as required by AutoDDG), and (2) openly accessible so we can extract samples for evaluation. We also manually screened citing papers to confirm that they made substantive use of the dataset.

4.2 Systems and Metrics

We use the released AutoDDG pipeline³ as a baseline and evaluate prompting variants of AutoDDG-Enhanced, which uses the PaperTrail module. We assess outputs using five complementary metrics spanning reference-based and reference-free evaluation. Reference-based metrics compare generations to the *repository-provided dataset descriptions* (i.e., descriptions on the dataset landing pages). Because these references are frequently incomplete and stylistically heterogeneous, we treat reference-based scores as informative but not definitive. We therefore emphasize reference-free metrics that directly assess research-facing summary quality, including coverage of key content, preference judgments, and LLM-based rubric scoring.

4.3 Reference-based Metrics

We report **ROUGE-1** and **ROUGE-L** to measure lexical overlap with the repository-provided descriptions, and **BERTScore** to measure semantic similarity. These metrics capture consistency with existing descriptions but are limited by the quality and completeness of the references.

4.4 Reference-free Metrics

We report (i) **coverage** under both **strict** and **lenient** matching across five coverage dimensions, (ii) **LLM-Preference** via pairwise comparisons, and (iii) **AutoDDG’s LLM-as-a-Judge** rubric scores for clarity, readability, and completeness. Because repository descriptions often omit the




Figure 3: Strict vs. lenient extraction on the baseline prompt. Lenient extraction is reported as a robustness check for indirectly stated high-level attributes; strict extraction is used as the primary coverage metric.

research-facing details we add, reference-based metrics cannot reliably capture these improvements; our reference-free metrics directly evaluate the presence and quality of this additional information, emphasizing dataset’s accessibility to researchers.

Coverage Score Strict/Lenient We create a measure of coverage by extracting a structured schema from each generated description and measuring which key attributes are present from our criteria. Concretely, we prompt an LLM to produce a JSON record with five dimensions: `basic_info` (weighted 20%), `data_characteristics` (weighted 30%), `provenance` (weighted 20%), `usage_context` (weighted 20%), and `quality_and_limitations` (weighted 10%). Each contains a set of fields outlining the extraction targets (e.g., `domain_or_field`, `size_or_scale`, `collection_method`, `typical_applications`, `known_limitations`).

We use **strict coverage score** as our primary metric because it rewards explicitly supported, auditable attributes. We additionally report a **lenient** variant as a sensitivity check that credits a small set of clearly implied high-level fields. In the **lenient** setting, the extractor may infer a small subset of high-level fields like domain/purpose and certain usage-context fields, we use this variant only to test robustness to indirect phrasing, while all other fields still require explicit evidence.

Figure 3 illustrates the expected effect of the lenient setting: as it increases overall scores and it increases scores primarily in *Basic Info* and *Usage*

³<https://github.com/VIDA-NYU/AutoDDG>

Context, where descriptions often express domain and intended use implicitly. We therefore treat lenient coverage as a secondary diagnostic, and use strict coverage score for primary comparisons.

Preference Ranking LLM We conduct a pairwise preference evaluation using an LLM-as-a-judge. For each dataset, the evaluator is given two descriptions side-by-side: the original AutoDDG output (A) and the augmented description produced by our system (B). The judge model is placed in the role of an experienced research scientist and instructed to compare the descriptions exclusively on their research utility, following a fixed rubric that scores completeness, accuracy, clarity, and actionability. The prompt emphasizes that longer descriptions should not be penalized when the added information is relevant, and that all decisions must be evidence-based. The evaluator outputs a structured JSON object containing: (1) a categorical preference (A, B, or Tie), (2) a 110 utility score for each description, and (3) a rationale that addresses each criterion. This setup ensures consistent, auditable comparisons and avoids stylistic or verbosity biases by requiring the model to justify its preference using the rubric.

5 Results

We first ask whether augmenting AutoDDG with PaperTrail improves the summarization generation with coverage of research-facing attributes. Table 1 reports strict and lenient overall coverage as well as LLM-as-judge scores and reference-based metrics for Vanilla AutoDDG and all PaperTrail prompting variants. Vanilla achieves an average strict coverage score of 0.397 ± 0.094 , whereas all context-augmented prompts perform at least as well, with a strict coverage score ranging from 0.39–0.46. Our best-performing prompt (RESEARCH_LONGV2) attains 0.46 ± 0.098 strict coverage score (+0.06 absolute over Vanilla).

Under the lenient extractor, which additionally credits clearly implied high-level fields, Vanilla reaches 0.37 ± 0.045 lenient coverage, while PaperTrail prompts span 0.40–0.46. The strongest variants (V1_REVISSED and RESEARCH_LONGV2) both reach 0.46 (within the standard deviation) lenient coverage. These trends indicate that PaperTrail consistently increases the amount of research-relevant information recoverable from the descriptions, and that the gains are not an artifact of the stricter extraction setting.




Figure 4: Best Prompt PaperTrail vs Vanilla AutoDDG

Beyond coverage, we evaluate whether PaperTrail affects the fluency and reference alignment of the generated descriptions.

LLM-as-a-Judge Metrics Results. Across prompts, completeness, readability, and conciseness remain stable relative to the Vanilla baseline. Most PaperTrail variants achieve completeness scores between 8.6–8.9 and readability scores between 8.6–9.0, indicating that adding research-derived context does not compromise clarity or narrative structure. Conciseness decreases mildly for some prompts (e.g., RESEARCH_LONGV2: 7.8 ± 0.632 vs. Vanilla’s 7.6 ± 1.174), which can be explained by descriptions becoming longer to incorporate provenance and usage details. However, the effect is modest.

Reference-Based Metrics Results. BERTScore and ROUGE-L exhibit only small fluctuations across prompting strategies. For instance, BERTScore improves from 0.825 (Vanilla) to 0.829 (RESEARCH_LONGV2), while consistently ROUGE-L varies between 0.01–0.02 across all prompts. This stability reflects the limitations of the repository-provided descriptions, because these references often omit the information PaperTrail strives to add on usage context, large gains in coverage do not necessarily correspond to higher lexical or semantic overlap with the reference text. Consequently, reference-based metrics provide only weak signals for our task, and their flat behavior encourages our use of strict coverage scoring as the primary evaluation measure.

As shown in Figure 6, across prompts, the largest improvements emerge in dimensions related to *provenance* and *usage context*. These dimensions are specifically the types of details that citation papers provide but that dataset summarization descriptions under report which indicates Papertrail succeeded at injecting new information on

Table 1: Overall performance by prompt (mean $\pm$ SD across 10 datasets). Best per column is bolded.

Prompt	Strict Cov.	Lenient Cov.	Completeness	Readability	Conciseness	BERT-F1	ROUGE-L
Vanilla	0.397 $\pm$ 0.094	0.367 $\pm$ 0.045	8.600 $\pm$ 0.516	8.700 $\pm$ 0.675	7.600 $\pm$ 1.174	0.825 $\pm$ 0.012	0.156 $\pm$ 0.018
V0_Original	0.401 $\pm$ 0.063	0.455 $\pm$ 0.070	8.900 $\pm$ 0.316	8.800 $\pm$ 0.422	7.900 $\pm$ 0.738	0.825 $\pm$ 0.010	0.149 $\pm$ 0.021
V1_Revised	0.435 $\pm$ 0.120	0.461 $\pm$ 0.104	8.900 $\pm$ 0.316	8.800 $\pm$ 0.632	7.800 $\pm$ 0.632	0.822 $\pm$ 0.008	0.138 $\pm$ 0.030
V2_Hybrid	0.444 $\pm$ 0.086	0.440 $\pm$ 0.080	8.800 $\pm$ 0.422	9.000 $\pm$ 0.000	7.800 $\pm$ 0.632	0.826 $\pm$ 0.008	0.154 $\pm$ 0.029
Structured_v1	0.390 $\pm$ 0.101	0.398 $\pm$ 0.067	8.900 $\pm$ 0.316	9.000 $\pm$ 0.000	8.100 $\pm$ 0.316	0.828 $\pm$ 0.009	0.160 $\pm$ 0.039
Research_longv1	0.447 $\pm$ 0.116	0.424 $\pm$ 0.109	8.700 $\pm$ 0.483	8.600 $\pm$ 0.843	7.400 $\pm$ 0.966	0.828 $\pm$ 0.007	0.152 $\pm$ 0.032
Research_shortv1	0.448 $\pm$ 0.099	0.444 $\pm$ 0.065	8.700 $\pm$ 0.483	9.000 $\pm$ 0.000	7.700 $\pm$ 0.949	0.826 $\pm$ 0.008	0.156 $\pm$ 0.035
Research_longv2	0.460 $\pm$ 0.098	0.459 $\pm$ 0.062	8.600 $\pm$ 0.966	8.600 $\pm$ 0.843	7.800 $\pm$ 0.632	0.829 $\pm$ 0.008	0.154 $\pm$ 0.038




Figure 5: Strict Coverage Dimension Scores by Prompt Types




Figure 6: Strict Coverage Dimension Scores by Prompt Types

research-focused context.

Despite substantial gains in strict coverage particularly for prompts like RESEARCH_LONGV2, the LLM preference evaluation shows a more mixed picture as seen in Figure 5. Across prompt types, the augmented descriptions were not consistently preferred over the baseline AutoDDG outputs, and in all cases the baseline description remained the majority winner. This reflects a limitation in automatic description generation, where adding more contextual information improves objective coverage but may reduce the perceived clarity as it can more variability in the structure of the descriptions.

We selected RESEARCH_LONGV2 as our default PaperTrail prompt for all subsequent analy-

ses as it achieves the highest strict overall coverage (0.46) among all prompts. In addition, RESEARCH_LONGV2 also performs strongly on lenient coverage and on LLM-as-a-judge metrics (completeness, readability, conciseness), showing no degradation in fluency or coherence. While some prompts slightly outperform it on individual secondary metrics, RESEARCH_LONGV2 is the only variant that performs robustly across all evaluation criteria without sacrificing strict coverage. We therefore adopt it as the representative PaperTrail configuration in the remainder of our experiments.

In Figure 4 we show the comparacment of Vanilla AutoDDG to our best PaperTrail prompt (RESEARCH_LONGV2) across all automatic metrics. As previously mentioned RESEARCH_LONGV2 achieves higher strict and lenient coverage, as well as slightly higher completeness and readability scores, indicating that the augmented descriptions capture more research-relevant information without hurting fluency. BERT-F1 and ROUGE-L remain essentially unchanged, reflecting that our additions go beyond what is present in the often-short repository descriptions. The only dimension where Vanilla slightly outperforms RESEARCH_LONGV2 is conciseness, consistent with the fact that the PaperTrail descriptions are longer in order to incorporate additional context from the citing papers.

6 Discussion, Limitations, Future Work

Our findings show that PaperTrail consistently improves coverage of research-relevant attributes especially provenance and usage context because these categories correspond directly to the kinds of information that scientific papers document but dataset repositories typically omit. Whereas AutoDDG relies solely on statistical and semantic profiling derived from the dataset itself, the cit-

ing papers provide complementary context such as data collection procedures, pre-processing decisions, typical analytical workflows, and known limitations observed in prior research. This additional information naturally boosts strict coverage in the provenance and usage dimensions, demonstrating that research literature can serve as a valuable external knowledge source for dataset summarization systems. Although PaperTrail improves objective coverage, its augmented descriptions are not consistently preferred by an LLM judge, with the baseline AutoDDG often winning head-to-head comparisons. This suggests that adding research-derived context can introduce structural variability or reduce perceived conciseness, even when the added information is relevant. PaperTrail makes descriptions more informative but also somewhat longer and less uniformly structured, which may affect perceived clarity. These results indicate that maximizing informational richness does not always align with maximizing user-perceived readability, and it is thus important to attain a balance between completeness and succinctness. Overall, the mixed preference results underscore that improvements in content quality do not automatically translate into improved user experience. This suggests that future work should explore mechanisms for better integrating research-derived insights into a stable structure template, as well as adaptive summarization strategies that can change the degree of added context to different depending on use cases.

References

- Kavita Ganesan. 2018. [Rouge 2.0: Updated and improved measures for evaluation of summarization tasks](#). *Preprint*, arXiv:1803.01937.
- Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Meng Wang, and Haofen Wang. 2024. [Retrieval-augmented generation for large language models: A survey](#). *Preprint*, arXiv:2312.10997.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2021. [Retrieval-augmented generation for knowledge-intensive nlp tasks](#). *Preprint*, arXiv:2005.11401.
- Haoxiang Zhang, Yurong Liu, Wei-Lun Hung, Aécio Santos, and Juliana Freire. 2025. [Autoddg: Automated dataset description generation using large language models](#). *Preprint*, arXiv:2502.01050.
- Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. 2020. [Bertscore: Evaluating text generation with bert](#). *Preprint*, arXiv:1904.09675.

Table 2: Datasets used in our experiment setup

Dataset	Domain	Zenodo
The FluPRINT database	Immunology / Vaccines	3222451
CODE-15%: a large scale annotated dataset of 12-lead ECGs	Clinical / ECG	4916206
T1GDUJA: Glucose dataset of a patient with type 1 diabetes mellitus	Clinical / Diabetes	11284018
A comprehensive database of expert-curated occurrences for the genus <i>Carex</i> L. (Cyperaceae)	Biodiversity / Occurrences	14998163
TetrapodTraits Database	Ecology / Traits	11303604
leafnp_data	Plant ecology / Biogeochemistry	11071944
Summary-level data from meta-analysis of fat distribution phenotypes in UK Biobank and GI-ANT	Genomics / GWAS	1251813
Coronavirus disease (COVID-19) case data – South Africa	Epidemiology / Public health	7651129
European database of processing factors for pesticides residues in food	Food safety / Pesticides	15363279
The Global Carbon Project’s fossil CO ₂ emissions dataset	Climate / Emissions	5569235